

FEATURE GEOMETRY FOR DATA-POOL SCREENING: UTILITY, STABILITY, AND LIMITS

Anonymous authors

Paper under double-blind review

ABSTRACT

Selecting source-data pools for a target task can require training and validating a predictor for many combinations of pools. Frozen features allow an earlier comparison, but retaining promising candidates is useful only if the final choice and the work saved justify screening. We evaluate a two-stage protocol that uses unlabeled source and target features to form a shortlist, then trains predictors on source labels and selects one using target validation. We compare target-weighted experimental design (A-opt) with second-moment matching (MMD) and target-blind diversity. The analysis explains the design score under a shared linear model and separates candidate omission from validation error. On equal-sample-cost DomainNet combinations, evaluating 50 A-opt or MMD candidates gives lower mean test Brier than evaluating 227 random candidates. Moment matching remains competitive across exploratory projection and readout changes. Domain-level matching also performs well after fixed-parameter repartitioning, while within-domain gains depend on construction and loss metric. Direct-selection gains on PACS and Office-Home are small, with limited improvement headroom. Independent timing shows savings with cached features. Adding measured feature-extraction costs preserves the logistic-regression saving but reverses it for ridge regression. These results assess screening through final selection quality and computation cost, rather than candidate retention alone.

1 INTRODUCTION

Researchers adapting a predictor to a target domain often have several source-data pools to choose from. The pools may differ in acquisition conditions, image style, or the examples they contain. To compare combinations, one can train a predictor on each and evaluate it on target validation data. The expense comes from repeating this process, even when each predictor is small. We ask whether an earlier comparison can reduce the number of combinations that receive supervised evaluation without discarding useful choices.

Pretrained encoders provide a common space for this comparison. An encoder maps each image to a feature vector. Keeping its parameters fixed allows those features to be reused across source combinations and classifier fits; CLIP and DINOv2 make this setting natural (Radford et al., 2021; Quab et al., 2024). A classifier trained on these fixed features is called a readout. Before fitting it, we can compare how source pools cover different feature directions and how that coverage relates to the target.

Broad coverage alone does not establish usefulness for the target. Effective rank summarizes how evenly variation is spread across directions by exponentiating the entropy of normalized singular values (Roy & Vetterli, 2007; Garrido et al., 2023). A high value may reflect variation rarely encountered at prediction time, and individually diverse pools may overlap. Source diversity and target relevance are therefore different questions. A score that uses target features can distinguish directions that a source-only spectral summary treats alike.

Geometry still leaves uncertainty about prediction: feature coverage does not by itself determine the relation between features and labels. We therefore use it to retain alternatives, then use labeled

validation to choose among them. This separates two questions. Does screening leave good candidates available, and does validation select a good predictor from that set? A useful ranking must ultimately be judged by the final choice and the work needed to obtain it.

We evaluate a two-stage data-pool screening protocol. Unlabeled source and target features first rank combinations. Predictors are then trained on the retained source combinations and selected using separate target-validation data. Target-weighted experimental design favors coverage that reduces uncertainty along target-relevant directions (Yu et al., 2006). Second-moment matching instead compares the source’s average directional energy with the target’s. We evaluate these scores against target-blind diversity, separate screening omissions from validation errors, and measure final loss across evaluation budgets and computation costs.

On equal-sample-cost DomainNet combinations, short geometric lists achieve lower mean test Brier, a probability-prediction loss, than longer random lists. Independent timing shows savings with cached features; adding measured feature work preserves the logistic-regression saving but reverses it for ridge regression. Source-domain controls and fixed-parameter repartitioning distinguish useful block information from domain-level matching, with outcomes depending on construction and metric. Projection and readout changes retain high average recall, whereas PACS and Office-Home leave little room to improve on direct-selection baselines. Together, these comparisons assess what target-aware geometry contributes to supervised pool selection and when its extra work is worthwhile.

2 RELATED WORK

2.1 OPTIMAL DESIGN AND DISTRIBUTION MATCHING

A-optimal design measures uncertainty through the trace of a covariance matrix, with applications to active linear regression (Fontaine et al., 2021). Transductive experimental design minimizes prediction variance over given target covariates (Yu et al., 2006); its regularized linear objective is the target-weighted trace used here. We apply this criterion to unions of fixed source pools and retain several unions for labeled validation. Distribution matching offers a different approach. MMD compares mean embeddings of source and target distributions (Gretton et al., 2012); a homogeneous quadratic kernel makes this comparison depend on their second moments. Moment alignment has also been used to learn representations for multi-source domain adaptation (Peng et al., 2019). Here the encoder stays fixed, allowing us to compare uncertainty reduction and moment matching as scoring rules on the same candidate combinations.

2.2 SELECTION UNDER A COMPUTATION BUDGET

Compute-constrained data selection accounts for selection work as well as subsequent training (Yin & Rush, 2025). Pool screening introduces an intermediate decision: how many combinations should receive supervised evaluation? A broad shortlist may retain good candidates but still require substantial fitting and validation. We compare both retention and the final selected model, then separately measure decision time and feature-preparation costs. This distinguishes fewer evaluations from a faster decision.

2.3 REPRESENTATION GEOMETRY AND DIVERSITY

RankMe uses effective rank to assess learned representations without downstream labels (Garrido et al., 2023). Determinantal point processes favor subsets with complementary directions (Kulesza & Taskar, 2012). Both motivate ways to compare source pools before fitting a task-specific model. Their geometric objectives differ, however, from preserving candidates with low target loss: variation may be redundant for a particular task, and directions important to that task may occupy only part of the representation. We examine merged effective rank and a specified subspace-DPP implementation as target-blind comparators.

3 GEOMETRY-GUIDED SCREENING

3.1 PROBLEM AND TWO-STAGE PROCEDURE

A pool is a fixed block of examples; a candidate is a feasible union of blocks. The study uses a candidate family and separate target data for screening, validation, and testing. Screening and supervised evaluation produce one selected predictor. A shortlist contains candidate combinations, not individual images.

Let $H_i \in \mathbb{R}^{n_i \times d}$ contain the features of pool i , using a fixed encoder and shared label-independent preprocessing. Its Gram matrix $G_i = H_i^\top H_i$ records feature-direction strengths and their relationships. For S in the candidate family $\mathcal{C}$, $G_S = \sum_{i \in S} G_i$ and $n_S = \sum_{i \in S} n_i$. Unlabeled target-selection features estimate the uncentered second moment $C_T = \mathbb{E}_T[hh^\top]$, which describes target directional energy.

The screen ranks $\mathcal{C}$ and retains its first s combinations as $\mathcal{Q}_s$. We then train predictors on their source labels and use a separate target-validation set to choose $\hat{S} \in \mathcal{Q}_s$. Target-test data evaluate this frozen choice and the candidates retained. The screening stage, denoted U0, sees source and target-selection features but no labels or validation/test data. The supervised stage, U2, reads source labels and target-validation data. Target-blind rules use only source information. This separation lets the screen exploit reusable features while reserving task-specific evidence for the final choice.

Each combination receives the same classifier and validation procedure; screening changes which combinations are evaluated, not the encoder. The equal-sample-cost study fixes training sample count, not elapsed time.

3.2 SHARED GEOMETRY AND SPECTRAL DIVERSITY

The singular value decomposition $H = U\Sigma V^\top$ separates feature directions from their strengths. For fixed nonnegative weights w_i , stacking the scaled source matrices gives

$$G_S = H_S^\top H_S = \sum_{i \in S} w_i G_i. \quad (1)$$

The eigenvalues of this sum are the squared singular values of the pooled features. It preserves the strength and cross-pool alignment of directions, provided preprocessing coordinates remain fixed. The experiments use unit weights.

Effective rank normalizes the singular values into a probability vector and exponentiates its entropy (Roy & Vetterli, 2007):

$$r_{\text{eff}}(H) = \exp\left(-\sum_j p_j \log p_j\right), \quad p_j = \frac{\sigma_j(H)}{\sum_k \sigma_k(H)}. \quad (2)$$

For a nonzero matrix, $0 \log 0 = 0$ and $1 \leq r_{\text{eff}}(H) \leq \text{rank}(H)$. Equal positive singular values attain the upper endpoint. Two singular values in ratio 9 : 1 give effective rank about 1.38, compared with two for equal strengths. We take square roots of Gram eigenvalues before normalization; normalizing the eigenvalues themselves would define a different score. Repeating an entire matrix leaves effective rank unchanged, so it measures spectral shape rather than sample evidence. It is a target-blind comparator, not an intermediate step required by the target-aware scores.

Marginal spectra and principal angles alone need not determine even this pooled statistic: the same strengths can be assigned differently to shared and unshared directions. Appendix A supplies an explicit counterexample and the conditions for a local spectral formula. We retain Gram operators when scoring combinations so that this directional information is available.

3.3 TARGET WEIGHTING AND MOMENT MATCHING

Target weighting asks which source combination reduces uncertainty where the target needs coverage. A shared linear model makes this question precise. Consider scalar Bayesian regression with $\theta \sim N(0, P_0)$, $P_0 \succ 0$, and $y = h^\top \theta + \epsilon$, where independent noise satisfies $\epsilon \sim N(0, \sigma^2)$ and

$\sigma^2 > 0$. Sources and target share θ and the noise model; target covariates reveal no further information about θ . Conditional on the source design, the posterior covariance and posterior-mean target risk are

$$P_S = (P_0^{-1} + \sigma^{-2}G_S)^{-1}, \quad R_T(S) = \sigma^2 + \text{tr}(C_T P_S). \quad (3)$$

The expectation averages posterior uncertainty and a new target observation. The derivation is recalled in Appendix A.

Equation (3) motivates the score $f_A(S) = \text{tr}(\hat{C}_T P_S)$, minimized with fixed hyperparameters. With an isotropic prior, this is the regularized target-prediction-variance criterion of transductive experimental design (Yu et al., 2006); we retain the name Target A-opt for continuity with the recorded experiments. Source observations reduce uncertainty in the directions they cover, and the target moment weights those reductions by how often the directions occur at prediction time. Thus a pool can be valuable even without high overall diversity if it covers directions that remain uncertain and matter to the target. Adding a positive-semidefinite source Gram cannot increase this score.

For the experimental choice $P_0^{-1} = I$ and $\sigma^2 = 1$, write $G_S = U \text{diag}(\lambda_j)U^\top$, including zero eigenvalues. Then

$$f_A(S) = \sum_j \frac{u_j^\top \hat{C}_T u_j}{1 + \lambda_j}. \quad (4)$$

The denominator reflects source coverage of direction u_j , while the numerator measures its importance to the target. Two combinations with the same eigenvalues can therefore receive different scores if their directions align differently with the target. Effective rank treats their spectra identically; target weighting distinguishes where their coverage is useful.

Second-moment matching asks a different question: how close is the source’s average directional energy to the target’s? It minimizes

$$f_M(S) = \left\| G_S/n_S - \hat{C}_T \right\|_F^2. \quad (5)$$

This is empirical squared MMD with kernel $k(h, h') = (h^\top h')^2$, including diagonal terms. It brings the source’s average directional energy closer to the target’s rather than weighting a posterior uncertainty matrix. Both criteria read the same features, so their comparison tests whether inverse-covariance weighting adds value beyond direct moment matching.

The shared conditional model supplies the risk interpretation of A-opt. Its assumption is substantive: if all features equal 1, source pools labeled constantly +1 and −1 produce predictors whose risk ordering reverses between two target worlds with opposite constant labels. The feature geometry is unchanged. In classification, we therefore evaluate the scores using trained ridge-softmax and logistic-regression readouts rather than identify Equation (3) with their losses.

3.4 SEPARATING RETENTION FROM FINAL SELECTION

The final choice can lose quality in two ways: screening discards a good candidate, or validation chooses a worse member of the retained set. Let $\ell(S)$ be target risk for a fixed trained predictor and $v(S)$ its validation loss. If $|v(S) - \ell(S)| \leq \varepsilon$ throughout $\mathcal{Q}_s$, the validation minimizer satisfies

$$\ell(\hat{S}) - \min_{S \in \mathcal{C}} \ell(S) \leq \underbrace{\min_{S \in \mathcal{Q}_s} \ell(S) - \min_{S \in \mathcal{C}} \ell(S)}_{\text{shortlist omission}} + 2\varepsilon. \quad (6)$$

Appendix A proves the bound. Its first term measures loss from excluding the best candidate; the second bounds validation error within the shortlist. We audit finite-test counterparts and top-ten retention. The deviation condition explains the two error sources; the experiments do not estimate it as a population guarantee.

3.5 DECISION BUDGETS AND COST ACCOUNTING

The full-search reference trains all candidates and chooses $S_{\text{full}} = \arg \min_{S \in \mathcal{C}} v(S)$. This is a deployable procedure. The test oracle, in contrast, is the candidate with lowest observed test loss and is used only for auditing. Let $b(S)$ be test Brier, the mean squared distance between predicted

class probabilities and the one-hot label; lower is better. We report $b(\hat{S}_s) - b(S_{\text{full}})$ and the fraction of repeated tasks exceeding a declared absolute tolerance. The gap can be negative because the validation and test rankings can differ.

With frozen features already available, a screening run costs its ranking work plus fitting and validating the retained candidates. Starting from images adds encoder initialization and feature extraction; target-aware scores additionally need target-selection features. These are different accounting scenarios. The break-even condition is that screening work be smaller than the supervised evaluations it replaces, after any additional feature work is included. Candidate counts measure one part of that condition, not elapsed time. In particular, the same matrix operations that make a linear score convenient can be a substantial part of fitting a ridge readout.

4 EVALUATION PROTOCOL

Controlled synthetic tasks test the risk model. PACS and Office-Home (Li et al., 2017; Venkateswara et al., 2017) then evaluate direct selection against a prespecified practical-effect threshold. An exhaustive headroom analysis follows those outcomes. The subsequent DomainNet study fixes sample cost and evaluates shortlists; its configuration was frozen before feature extraction and outcome evaluation. Test-sample resampling and the later projection/readout study are exploratory follow-ups. Appendix B records their configurations and result locations.

From the official cleaned DomainNet data (Peng et al., 2019), we select 100 class names by a fixed hash rule. Each of six domains contributes disjoint sets of 1,536 target-selection images, 1,024 target-validation images, and 3,072 anchor-pool images from official training data, plus 2,048 official test images. The manifest contains 46,080 images. Class metadata define the benchmark vocabulary but are hidden from screening.

CLIP ViT-B/32 (Radford et al., 2021) serves as a frozen construction anchor. Within each domain, its first principal-component scores define low, middle, and high strata, each supplying a block of 512 images. For a given target, we exclude its own blocks and enumerate $\binom{15}{3} = 455$ combinations, each with 1,536 source images. ResNet-50 and DINOv2-B/14 provide evaluation features (He et al., 2016; Oquab et al., 2024). All rules share row normalization, a fixed Rademacher projection to 128 dimensions, and post-projection row normalization. Holding sample count constant removes a direct source of training-cost differences between candidates.

Target A-opt uses $P_0^{-1} = I$ and $\sigma^2 = 1$. Comparators are second-moment MMD, target energy $\text{tr}(\hat{C}_T G_S / n_S)$, Bayesian D-opt, merged effective rank, subspace DPP, domain balance, and 20 seeded random rankings. They score the same candidate family. Domain balance additionally reads source-domain identifiers. Kernel definitions and tie rules appear in Appendix B.

The screen produces nested shortlists of sizes 10, 25, 50, 100, 227, with 227 designated primary. In the original protocol, U2 independently evaluates each method’s primary shortlist using ridge regression with regularization 1, no intercept, and a fixed softmax temperature. Validation Brier selects the final combination, breaking ties lexicographically. Smaller lists reuse that method’s nested evaluations; there is no cross-method cache. The test stage verifies upstream hashes before evaluating every combination to establish an exhaustive audit reference.

The follow-up adds four projection seeds and a multinomial logistic readout with fixed regularization. It keeps the candidate members and data roles unchanged. Screen, validation, and test remain separate processes, but the follow-up reuses each combination’s fitted model across methods to audit their shortlists. Its runtime is therefore the cost of a cached research audit rather than separate deployments. Readout settings and provenance checks are detailed in Appendix B.

We measure how many test top-ten candidates survive screening and how far the validation choice lies from the test oracle S^* . Recall measures candidate retention, not classification accuracy; regret measures loss relative to the oracle:

$$\text{Recall}_{10}(\mathcal{Q}_s) = \frac{|\mathcal{Q}_s \cap \text{Top}_{10}(b)|}{10}, \quad \text{Regret}(\hat{S}) = \frac{b(\hat{S}) - b(S^*)}{|b(S^*)|}. \quad (7)$$

The original acceptance criteria require at least 50% removal, mean recall of at least 90%, and mean normalized regret at most 5%. Brier is primary. Follow-up negative log-likelihood (NLL) and




Figure 1: Mean final test Brier minus the exhaustive-validation choice is plotted against the number of candidates evaluated. The two panels use different vertical scales. The zero line is a validation-selected model, not the test oracle. Curves are post-hoc summaries of the four new projection seeds; the horizontal axis counts fits, not measured runtime.

classification-error audits use their corresponding validation metric to select candidates, with the metric fixed before test access.

We average encoder results within each dataset–target-domain unit and random rankings within each task. PACS/Office-Home contain eight such units and DomainNet six. Follow-up seeds are also averaged within domains. Original bootstrap intervals resample target units; the test-sample analysis fixes trained models and resamples examples within each target. Seeds, encoders, and overlapping combinations are repeated measurements, not additional independent datasets.

5 EMPIRICAL FINDINGS

We first examine final choices and computation cost on DomainNet, then ask whether the advantage comes from domain composition or from choosing blocks within a domain. Retention and readout analyses explain which alternatives remain available. The earlier shared-model and direct-selection studies provide context for interpreting these results; they are not additional equal-cost replications.

5.1 FINAL CHOICES ACROSS EVALUATION BUDGETS

Figure 1 reanalyzes the four new projection seeds using final test loss rather than top-ten membership. The reference fits all candidates and selects by the same validation metric. At small budgets, target-aware lists have lower mean Brier than equally sized random lists. Their curves approach the exhaustive-validation reference well before all candidates are evaluated. The practical comparison is therefore between decision procedures at a specified loss tolerance, not simply between ranking correlations.

For logistic regression, evaluating 50 A-opt or MMD candidates gives lower mean Brier than evaluating 227 random candidates. Both comparisons favor screening in five of six domain averages. The paired mean differences are -0.001518 for A-opt and -0.001664 for MMD; descriptive domain-bootstrap ranges are $[-0.003517, 0.000023]$ and $[-0.003720, -0.000119]$. The domains share source pools, so these are not confidence statements about independent datasets.

The task-level exceedance rates show what the means omit. At size 50, each geometric method exceeds an absolute Brier gap of 0.005 to exhaustive validation in one of 48 repeated encoder–seed–domain tasks, or 2.08%. Random screening at size 227 exceeds it in 15.52% after the prescribed averaging. Appendix Table 5 reports the final losses. This tolerance is one point in an exploratory sensitivity grid, not a new success gate. The comparison describes both average quality and the frequency of larger gaps within the evaluated tasks.

5.2 MEASURED DECISION AND FEATURE COSTS

Table 1 tests whether evaluating fewer candidates saves time. Each method scores, fits, and validates independently; no fitted model is shared across methods. Quality and time come from the same runs. Unlike the preceding four-seed analysis, this profile fixes one projection and uses three executions per encoder–target–readout task. Appendix Figure 4 shows all measured budgets.

With projected features already in memory, A-opt and MMD at size 50 take about one tenth of exhaustive logistic-regression time. Adding the measured image-feature work leaves their accounted cost at about 56% of exhaustive validation. For ridge, extra target-selection features cost more than the cheap fits screening avoids. Thus the same reduction in candidate count has different consequences for the two readouts. Image totals add separately measured components; they exclude pool construction, downloads, cache serialization, and test auditing. Disk-cache loading and projection are reported separately in Appendix Table 7.

Table 1: Matched final Brier and mean accounted time in seconds are reported for the two readouts. Resident time includes scoring, fitting, and validation from in-memory projected features. Images adds the separately measured feature work for the fixed pool inventory; it is an additive profile, not a complete pipeline replay. All means use the same timed choices.

Readout	Screen	Fits	Brier	Resident (s)	Images (s)
Ridge	A-opt	50	0.986655	0.53	39.59
Ridge	MMD	50	0.986643	0.21	39.27
Ridge	Random	227	0.986653	0.75	35.30
Ridge	Full validation	455	0.986639	1.48	36.03
Logistic	A-opt	50	0.863170	4.49	43.54
Logistic	MMD	50	0.863110	4.31	43.37
Logistic	Random	227	0.864241	21.23	55.77
Logistic	Full validation	455	0.862992	42.83	77.38

5.3 CANDIDATE CONSTRUCTION AND SOURCE-DOMAIN CONTROLS

We next fix the projection and readout parameters and replace the CLIP-stratified blocks with a label- and feature-independent hash partition of the same source images. Mean matching and domain-level moment matching provide simpler controls. The latter pools all source blocks within each domain before scoring domain mixtures, so it cannot distinguish blocks with identical source-domain composition. Appendix Table 6 reports this exploratory challenge; randomized tie handling is reported separately from the frozen lexicographic protocol.

The new construction preserves the broad Brier advantage over random screening, but changes which geometric details help. At size 50, domain-level matching and MMD have similar mean final Brier after repartitioning. Good global shortlists can therefore be obtained here using domain-level moments without distinguishing individual blocks within a domain. Mean matching is competitive on the original blocks but less effective after repartitioning.

To compare blocks while holding domain composition fixed, we group candidates that contain the same number of blocks from each source domain. We retain the floor of half the candidates in each non-singleton group. With the original blocks and logistic readout, A-opt and MMD retain the empirical Brier optimum in about 77% and 76% of groups, compared with 45% for random selection. After repartitioning, these rates fall to about 56% and 55%; for classification error, both fall below random retention. These are within-group optimum-retention rates, not global top-ten recall. They show useful block information on the original construction, with a benefit that depends on construction and metric.

5.4 RETAINING PROMISING EQUAL-COST COMBINATIONS

The original DomainNet study asks which candidates remain available after screening (Table 2). A-opt retains nearly all test top-ten candidates while discarding approximately half the combinations. Every primary A-opt and MMD shortlist contains a test Brier oracle; their remaining selected regret

Table 2: The original DomainNet study retains $s = 227$ of 455 combinations. Means average encoders within six target domains. Recall concerns the test Brier top ten; regret is relative to oracle Brier. Spearman measures agreement between the complete screening and test-loss rankings.

Screen	Recall (%)	Regret (%)	Spearman
Target A-opt	98.333	0.001311	0.7443
Second-moment MMD	98.333	0.001311	0.6902
Bayesian D-opt	82.500	0.002658	0.4996
Merged effective rank	81.667	0.002909	0.4951
Random (20 repeats)	48.958	0.003899	0.0016
Target energy	46.667	0.009347	-0.0399
DPP subspace	27.500	0.009067	-0.0633
Domain balance	25.000	0.012083	0.0681

comes from validation choosing another member. This illustrates Equation (6): retaining the best candidate and selecting it are separate outcomes.

MMD reaches the same primary recall and selected regret as A-opt, and validation chooses the same final combination for both in every encoder–target task. At smaller shortlist sizes, MMD has higher mean recall (Appendix Figure 2). A-opt’s stronger full-ranking correlation therefore does not translate into better retention at those operating points. The common advantage over the tested target-blind rules supports using target-conditioned second-order information; direct moment matching is already effective here.

5.5 PROJECTION AND READOUT SENSITIVITY

The original ridge-softmax losses differ little: every candidate lies within the 5% regret gate, so that gate does not distinguish screens. The exploratory follow-up asks whether high recall persists with four new projections and logistic regression, keeping source membership and sample cost fixed. Appendix Table 3 reports the new-seed averages; the original seed is an implementation control.

A-opt retains 98.33% of the Brier top ten on average under each readout, with MMD slightly higher. The mean relative best-to-worst Brier span increases from 0.248% under ridge-softmax to 14.65% under logistic regression. High recall thus persists in a setting with more separation between candidates, rather than appearing only when their Brier losses are almost equal. The two target-aware scores also retain most NLL- and error-ranked candidates on average, with MMD higher in these comparisons.

At the primary shortlist size, both methods retain a Brier test oracle in every new-seed task and have the same mean selected regret. Metric-specific outcomes vary: one A-opt task has only 40% error-ranked recall, and validation misses a retained test-error oracle. Appendix C gives the task details and full comparator table.

5.6 TARGET WEIGHTING UNDER THE SHARED RISK MODEL

The earlier synthetic study tests target weighting when its risk model holds. It varies feature dimension, candidate count, source budget, target sample size, and target rank, using 20 generation seeds for 216 design settings. Under the shared conditional model, A-opt reduces modeled risk by 4.23% relative to the strongest target-blind effective-rank comparator and by 14.85% relative to MMD. The respective descriptive 95% intervals are [3.86, 4.62]% and [13.99, 15.70]%, resampling design settings after averaging seeds. These comparisons concern modeled linear risk, rather than the classification losses above.

A separate source-conditional-shift sweep changes the regression mechanism. Its mean normalized prediction regret increases from 0.211 without shift to 3.143 at the largest tested shift. This endpoint measures realized prediction rather than posterior risk. Together, the two studies distinguish the advantage of covering target-relevant directions from the effect of transferring labels across different conditional mechanisms.

5.7 DIRECT SELECTION AND AVAILABLE IMPROVEMENT HEADROOM

PACS and Office-Home test whether the score can directly choose a better source combination. At block budget $K = 3$, A-opt improves mean within-unit relative Brier by 0.351% over DPP and 0.273% over MMD. Both fall below the prespecified 2% practical-effect threshold. The paired difference interval crosses zero against DPP and is negative against MMD; detailed values are in Appendix C. These comparisons hold block count, rather than source sample count, fixed.

The post-hoc audit evaluates every feasible combination to measure improvement headroom: the loss reduction available by choosing the test-best candidate rather than the reference. Appendix Figure 3 shows only 0.781% mean headroom over DPP. Most target units leave little room for a selector to improve the average loss substantially. This distinguishes direct predictive gains from the screening question: whether fewer combinations can be evaluated while retaining useful choices.

6 DISCUSSION

6.1 WHAT MAKES SCREENING USEFUL?

The experiments assess screening through three connected outcomes: which candidates survive, which predictor validation selects, and how much work the decision requires. High recall alone answers only the first question. The MMD and domain-matching results also make simple target comparisons essential baselines for more detailed block geometry.

The operating point should reflect the target metric, acceptable loss gap, and measured cost, rather than retention alone. Compression introduces another requirement: the score must still distinguish candidates. In Appendix C, valid sketch intervals remain too wide to resolve choices at a useful byte budget.

6.2 LIMITATIONS

The equal-sample-cost evidence comes from two constructions on the same DomainNet images. Its target domains share source pools, and the added seeds and readouts reuse the same images. They measure sensitivity within this benchmark rather than replication on independent data. Changes from PACS/Office-Home also include the dataset, projection, sample budget, and candidate construction, so they do not isolate candidate heterogeneity as a cause.

The Bayesian derivation concerns shared linear regression; the classification results establish the usefulness of its score empirically. Test-example resampling conditions on fitted models and validation choices. It does not measure uncertainty from training or validation selection. Absolute task quality also depends on the candidate family, as the supplementary adapter study illustrates. Image costs are measured components, not a replay of a complete deployment: pool construction and cache serialization are excluded, and the feature profile uses one pass on one hardware setup. Independent equal-cost datasets and controlled validation budgets remain important, as does checking overlap with encoder pretraining data.

7 CONCLUSION

We evaluated target-aware geometry as a screening step before supervised data-pool selection. The analysis connects experimental design to shared linear risk and separates candidate omission from validation error. On equal-sample-cost DomainNet combinations, evaluating 50 A-opt or MMD candidates gives lower mean test Brier than evaluating 227 random candidates. Domain controls show useful but construction-dependent block information, with MMD competitive in the DomainNet screening comparisons. Independent timing supports savings with cached features; adding measured feature costs preserves the logistic-regression saving but reverses it for ridge. The resulting evidence supports judging a screen by its final choice, acceptable loss gap, and measured cost together.

AI USE STATEMENT

Generative AI tools assisted conceptual reformulation, mathematical claim and proof drafting, experiment design, software implementation, interpretation of results, literature search, manuscript writing, figure scripting, and Chinese translation. Numerical findings in this draft are taken from recorded experiment artifacts rather than generated observations. AI assistance is not independent validation of those artifacts or of the mathematical claims. Final verification and responsibility for a submitted version remain with the human authors.

REPRODUCIBILITY STATEMENT

Proof details appear in Appendix A; preprocessing, access rules, candidate construction, metrics, and result paths appear in Section 4 and Appendix B. The code repository is <https://anonymous.4open.science/r/collapsemodel-1DB6>. Outcome tables and manifests support aggregate reproduction; large feature caches are identified by hashes rather than duplicated in Git. Exact hardware-independent bitwise feature reproduction is not claimed.

ETHICS STATEMENT

The study uses existing benchmark images and pretrained encoders, with no new participant recruitment. Public availability does not remove image licensing, privacy, or representational-bias concerns. We do not claim that a diverse feature spectrum or a high shortlist recall makes a selected pool socially representative, fair, or private. The candidate taxonomy, pretraining overlap uncertainty, and restriction to visual classification limit downstream deployment claims. Dataset and model access remain subject to their original terms.

REFERENCES

- Koenraad M. R. Audenaert. A sharp continuity estimate for the von Neumann entropy. *Journal of Physics A: Mathematical and Theoretical*, 40(28):8127–8136, 2007. doi: 10.1088/1751-8113/40/28/S18.
- Xavier Fontaine, Pierre Perrault, Michal Valko, and Vianney Perchet. Online A-optimal design and active linear regression. In *International Conference on Machine Learning*, pp. 3374–3383, 2021. URL <https://dblp.org/rec/conf/icml/FontainePVP21>.
- Quentin Garrido, Randall Balestriero, Laurent Najman, and Yann LeCun. RankMe: Assessing the downstream performance of pretrained self-supervised representations by their rank. In *International Conference on Machine Learning*, pp. 10929–10974, 2023. URL <https://dblp.org/rec/conf/icml/GarridoBNL23>.
- Arthur Gretton, Karsten M. Borgwardt, Malte J. Rasch, Bernhard Schölkopf, and Alexander Smola. A kernel two-sample test. *Journal of Machine Learning Research*, 13(25):723–773, 2012. URL <https://dblp.org/rec/journals/jmlr/GrettonBRSS12>.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *IEEE Conference on Computer Vision and Pattern Recognition*, pp. 770–778, 2016. URL <https://dblp.org/rec/conf/cvpr/HeZRS16>.
- Alex Kulesza and Ben Taskar. Determinantal point processes for machine learning. *Foundations and Trends in Machine Learning*, 5(2–3):123–286, 2012. doi: 10.1561/22000000044. URL <https://dblp.org/rec/journals/ftml/KuleszaT12>.
- Da Li, Yongxin Yang, Yi-Zhe Song, and Timothy Hospedales. Deeper, broader and artier domain generalization. In *IEEE International Conference on Computer Vision*, 2017. URL <https://dblp.org/rec/conf/iccv/LiYSH17>.
- Maxime Oquab, Timothée Darcet, Théo Moutakanni, Huy Vo, Marc Szafraniec, Vasil Khilodov, Pierre Fernandez, Daniel Haziza, Francisco Massa, Alaaeldin El-Nouby, Mahmoud Assran, Nicolas Ballas, Wojciech Galuba, Russell Howes, Po-Yao Huang, Shang-Wen Li, Ishan Misra, Michael Rabbat, Vasu Sharma, Gabriel Synnaeve, Hu Xu, Hervé Jégou, Julien

- Mairal, Patrick Labatut, Armand Joulin, and Piotr Bojanowski. DINOv2: Learning robust visual features without supervision. *Transactions on Machine Learning Research*, 2024. URL <https://dblp.org/rec/journals/tmlr/OquabDMVSKFHMEA24>.
- Xingchao Peng, Qinxun Bai, Xide Xia, Zijun Huang, Kate Saenko, and Bo Wang. Moment matching for multi-source domain adaptation. In *IEEE/CVF International Conference on Computer Vision*, pp. 1406–1415, 2019. URL <https://dblp.org/rec/conf/iccv/PengBXHSW19>.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning*, pp. 8748–8763, 2021. URL <https://dblp.org/rec/conf/icml/RadfordKHRGASAM21>.
- Olivier Roy and Martin Vetterli. The effective rank: A measure of effective dimensionality. In *15th European Signal Processing Conference*, pp. 606–610, 2007. URL <https://dblp.org/rec/conf/eusipco/RoyV07>.
- Hemanth Venkateswara, Jose Eusebio, Shayok Chakraborty, and Sethuraman Panchanathan. Deep hashing network for unsupervised domain adaptation. In *IEEE Conference on Computer Vision and Pattern Recognition*, pp. 5018–5027, 2017. URL <https://dblp.org/rec/conf/cvpr/VenkateswaraECP17>.
- Junjie Oscar Yin and Alexander Rush. Compute-constrained data selection. In *International Conference on Learning Representations*, 2025. URL <https://dblp.org/rec/conf/iclr/YinR25>.
- Kai Yu, Jinbo Bi, and Volker Tresp. Active learning via transductive experimental design. In *Proceedings of the 23rd International Conference on Machine Learning*, pp. 1081–1088, 2006. doi: 10.1145/1143844.1143980. URL <https://dblp.org/rec/conf/icml/YuBT06>.

A PROOFS AND SCOPE OF THE MATHEMATICAL CLAIMS

Proposition 1 (Directional assignment matters). *There exist row-normalized pairs (A, B_1) and (A, B_2) with identical marginal singular spectra and principal angles between their row spaces, but different effective ranks after concatenation.*

Proof. Take A with four rows $e_1^\top$ and one row $e_2^\top$. Let B_1 have four rows $e_1^\top$ and one $e_3^\top$, and let B_2 have one row $e_1^\top$ and four $e_3^\top$. All three nonzero singular spectra are $(2, 1)$; both row-space pairs have principal angles $(0, \pi/2)$. Their merged Gram spectra are $(8, 1, 1)$ and $(5, 4, 1)$, giving effective ranks approximately 2.6260 and 2.8495. The shared direction receives different energy despite the identical marginal summaries. $\square$

Let u, v be unit directions with $u^\top v = c \in [0, 1]$. In the orthonormal basis u, z , where $v = cu + \sqrt{1 - c^2}z$, the operator $a^2 uu^\top + b^2 vv^\top$ is

$$\begin{pmatrix} a^2 + b^2 c^2 & b^2 c \sqrt{1 - c^2} \\ b^2 c \sqrt{1 - c^2} & b^2 (1 - c^2) \end{pmatrix}. \quad (8)$$

Its trace is $a^2 + b^2$ and determinant $a^2 b^2 (1 - c^2)$, so its eigenvalues are

$$\lambda_{\pm} = \frac{a^2 + b^2 \pm \sqrt{(a^2 - b^2)^2 + 4a^2 b^2 c^2}}{2}. \quad (9)$$

At $c = 1$, the second branch is zero. These pairs give a full spectrum when the associated blocks are mutually orthogonal and invariant for both source Grams; unpaired directions contribute their own eigenvalues. This compatibility condition matters because principal vectors need not be singular vectors.

For fixed H_S and Gaussian likelihood, completing the square in the prior times likelihood gives posterior precision $P_0^{-1} + \sigma^{-2} H_S^\top H_S$ and mean $\mu_S = P_S \sigma^{-2} H_S^\top y_S$. For an independent target observation $y = h^\top \theta + \epsilon$, conditional on the training data and h ,

$$\mathbb{E}[(y - h^\top \mu_S)^2 \mid H_S, y_S, h] = \sigma^2 + h^\top P_S h. \quad (10)$$

Integrating h gives Equation (3). We assume the target covariate distribution does not reveal additional information about θ beyond this model. If a source has a different regression function, or labels are multiclass with softmax readout, this derivation no longer identifies the evaluated loss. The proof does not bound this misspecification error. Replacing C_T by a finite-sample estimate adds another discrepancy not removed by the identity.

If $G \succeq \tilde{G} \succeq 0$, inverse order implies $(P_0^{-1} + \sigma^{-2} G)^{-1} \preceq (P_0^{-1} + \sigma^{-2} \tilde{G})^{-1}$. Taking a trace against a PSD target moment preserves this order. More generally, suppose the sum of omitted sketch tails satisfies $0 \preceq G - \tilde{G} \preceq \delta I$. Define $B = P_0^{-1} + \sigma^{-2} \tilde{G}$. Then

$$\text{tr}(\hat{C}_T (B + \sigma^{-2} \delta I)^{-1}) \leq f_A(G) \leq \text{tr}(\hat{C}_T B^{-1}). \quad (11)$$

One admissible δ is the sum of valid spectral-norm bounds on each PSD tail. A candidate’s proxy ordering is certified when its upper bound is below every competing lower bound. Thus interval width, as well as coverage, determines whether a sketch can resolve the next choice. The archived sweep evaluates this type of proxy comparison with its documented byte accounting.

Let S_Q minimize ℓ over $\mathcal{Q}_s$. Uniform deviation on that shortlist gives

$$\ell(\hat{S}) \leq v(\hat{S}) + \varepsilon \leq v(S_Q) + \varepsilon \leq \ell(S_Q) + 2\varepsilon. \quad (12)$$

Subtracting $\min_{S \in \mathcal{C}} \ell(S)$ proves Equation (6). No probability is asserted by this algebra. For population generalization, one must additionally justify the deviation event using the actual dependence and selection scheme. Exhaustive finite-test scores alone do not supply that justification.

The feature map $\phi(h) = hh^\top$, equipped with Frobenius inner product, satisfies $\langle \phi(h), \phi(h') \rangle_F = (h^\top h')^2$. The squared distance between source and target empirical feature means is therefore Equation (5). This is the biased empirical squared MMD (a V-statistic), not the diagonal-excluded unbiased estimator. It depends on uncentered second moments and cannot distinguish all distributions sharing those moments. We use it as a computationally simple same-information comparator, not a complete discrepancy measure.

B IMPLEMENTATION AND AUDIT DETAILS

The original DomainNet study fits a ridge readout and transforms its scores into class probabilities. Its weights and Brier loss are

$$W_S = (I + G_S)^{-1} \sum_{i \in S} H_i^\top Y_i, \quad p(x) = \text{softmax}(h(x)^\top W_S),$$

$$b(S) = \frac{1}{n_T} \sum_t \|p(x_t) - e_{y_t}\|_2^2. \quad (13)$$

Y_i is the one-hot source-label matrix over the fixed 100-class vocabulary. There is no intercept or temperature tuning. Linear algebra uses float64 after loading float32 features. For effective rank obtained through a Gram eigendecomposition, negative roundoff eigenvalues are clipped to zero, square roots form singular values, and only $\sigma_j > \sigma_{\max} \sqrt{8\epsilon_{64} \max(n_S, d, 1)}$ are retained. The zero matrix returns rank zero in code. This tolerance concerns numerical rank; the entropy definition itself is unchanged. The shared projection uses seed 20260911 and $d = 128$; no outcome-dependent dimension choice is used within the frozen DomainNet protocol.

Bayesian D-opt maximizes $\log \det(I + G_S)$. Subspace DPP takes each block’s top 16 Gram eigenvectors U_i , defines $K_{ij} = \|U_i^\top U_j\|_F^2 / 16$ with $K_{ii} = 1$, numerically projects the matrix to a PSD correlation matrix, and maximizes $\log \det(K_{S,S} + 10^{-8}I)$. Its kernel describes block subspaces, not individual-image features. Domain balance maximizes the number of represented domains minus 10^{-3} times the sum of squared domain block counts. Deterministic scores break ties lexicographically by combination identifier. These choices are fixed comparators; poor performance of this DPP construction is not a result about all determinantal selection methods.

DomainNet records image content hashes, sample identifiers, official split membership, role assignments, and class vocabulary. Archive and split checksums are validated before feature extraction. Disjoint role membership and duplicate-content checks address overlap within the materialized benchmark; they do not rule out semantic near-duplicates or pretraining contamination. All screen methods use the same source blocks, projection, target-selection rows, and candidate family. Feature extraction ignores labels, while benchmark materialization uses taxonomy metadata and stores labels for later authorized stages.

Screen, validation, and audit manifests bind configuration and upstream artifact hashes. Stage-specific loaders enforce label and split permissions. In the original protocol, validation evaluates each method or random repeat separately. The test audit covers all 6×455 combinations per encoder after choices are frozen; its exhaustive oracle serves as a research reference. The follow-up caching policy is described below.

The CUDA reconstruction retains the original sample identifiers, candidate membership, preprocessing, and hyperparameters. All twelve Brier top-ten sets agree with the original study. Per-example loss means reproduce the reconstructed aggregate audit within 1.78×10^{-15} . For each target domain, 2,000 paired bootstrap draws resample its 2,048 test examples, sharing weights across combinations, encoders, and four losses. Fitted models and validation choices stay fixed; combination identifiers resolve ties. These draws measure sensitivity to test-example sampling. The original Brier top ten overlaps with NLL by 90.83%, raw squared loss by 30.83%, and classification error by 29.17% on average.

Seeds 20260917, 20260918, 20260919, and 20260920 generate new 128-dimensional Rademacher projections with the same pre/post row normalization. Seed 20260911 is an implementation control, excluded from new-seed means. Logistic regression uses L2 regularization, $C = 1$, no intercept, LBFGS, tolerance 10^{-6} , and at most 1,000 iterations. A convergence warning aborts a task; no such warning occurred, and the largest iteration count was 24. The class vocabulary is fixed, with zero probability assigned to source-absent classes.

All methods retain their screening rules, candidate combinations, and five shortlist sizes. Separate processes screen, fit and validate, then test frozen models. The audit checks source, configuration, lineage, model, and selection hashes before reading test features. The follow-up fits every combination once per readout and reuses these fits across methods. This cached exhaustive evaluation supports counterfactual comparisons of shortlists. Brier, NLL, and error use their own validation

losses for final selection. Original-seed ridge reproduces 5,460 reconstructed test combinations with maximum loss difference 1.78×10^{-15} and all 1,620 Brier validation choices exactly.

The two encoders and five seeds produce 54,600 fits across both readouts. New-seed summaries first average random repeats within a task, then encoders and seeds within each target domain, and finally the six domains. Each readout contributes 48 repeated encoder–seed–domain measurements. They remain six domain units sharing the same source pools.

The archive retains historical experiment IDs to identify results without confusing them with the paper’s scientific questions. The following paths are relative to the anonymous code repository:

- Shared-model and shift studies (E1): `results/target_conditioned/e1_synthetic/expanded_v3_controlled_rng/`.
- PACS/Office-Home comparison (E2): `results/target_conditioned/e2_dataset_selection/method_v1/`.
- Post-hoc headroom (E2a): `results/target_conditioned/e2_dataset_selection/oracle_headroom_v1/`.
- Equal-cost shortlist (E2b): `results/target_conditioned/e2b_fixed_cost_shortlist/domainnet_v1/`.
- Risk sketch sweep (E5): `results/target_conditioned/e5_sketch_certificate/expanded_v3_controlled_rng/`.
- Reconstruction and test resampling (R1/R2): `results/target_conditioned/revision_20260916/`.
- Projection and readout follow-up (R3/R4): `results/target_conditioned/revision_20260917/`.

The DomainNet data manifest begins `1ba25059df70be4d`, candidate-set ID `545fe86c8c9e8304`, and summary ID `2d473cb9eb977e85`; full SHA-256 identifiers are in the artifacts. The original configuration is `code/configs/target_conditioned_e2b/domainnet_v1.json`, implemented in `scripts/run_target_conditioned_e2b_shortlist.py`. The follow-up uses `code/configs/target_conditioned_e2b/readout_projection_v1.json` and `scripts/run_e2b_readout_projection.py`. Its summary tables are in `revision_results/e2b_readout_projection_v1/summary/` under the R3/R4 directory. The manuscript’s `build_assets.py` regenerates tables, figures, numeric macros, and bilingual abstracts from recorded sources, logging input hashes in `generated/asset_manifest.json`.

The three original feature matrices are not duplicated in Git. Metadata record their exact file and checkpoint hashes, shapes, preprocessing, extraction code revision, and runtime. Re-extraction requires access to the original public data and model weights; bitwise identity across different hardware is not assumed. The repository contains manifests and sufficient outcome tables to reproduce reported aggregate results without those matrices.

C SUPPLEMENTARY RESULTS

At $K = 3$ on PACS/Office-Home, mean Brier is 0.826769 for A-opt, 0.829162 for DPP, and 0.828719 for MMD. The paired A-opt-minus-baseline intervals are $[-0.007449, 0.000731]$ and $[-0.003392, -0.000712]$, respectively. The post-hoc audit covers all 220 combinations per task. Oracle headroom over DPP has descriptive 95% interval $[0.162, 1.617]\%$; only PACS photo exceeds the original 2% threshold at the target-unit level. The oracle improves on the post-hoc test-best envelope of parent methods by 0.087% on average. This envelope is an outcome-based reference, not a selectable baseline.

Table 4 gives the complete Brier recall comparison at the primary size. A-opt and MMD have the same mean relative selected regret: 0.001006% for ridge and 0.103% for logistic regression. All their new-seed primary Brier shortlists contain an empirical oracle.




Figure 2: Mean test Brier top-ten recall is shown across shortlist sizes in the original DomainNet study. The dotted line marks the prespecified recall threshold. All curves use the same candidate combinations and feature access appropriate to each rule.

Table 3: The exploratory DomainNet follow-up evaluates shortlists of size $s = 227$. Test top-ten recall (%) is averaged over two encoders and four new projection seeds within each target domain, then over six domains. Ridge denotes ridge-softmax. NLL and error use matching validation losses for final selection.

Readout	Metric	A-opt	MMD	Random
Ridge	Brier	98.33	99.38	50.12
Ridge	NLL	98.54	99.38	49.86
Ridge	Error	96.25	97.50	48.63
Logistic	Brier	98.33	98.96	48.97
Logistic	NLL	98.54	99.79	48.61
Logistic	Error	97.29	97.92	49.04

The lowest A-opt classification-error recall occurs for DINOv2, Quickdraw, logistic regression, and projection seed 20260918. The test-error oracle is 85.9863%, compared with 87.6465% for the validation choice. Its oracle remains in the shortlist despite the 40% top-ten recall. Thus retaining the best member and retaining the best ten answer different questions, and both are distinct from the absolute quality of the available models.

The four-encoder, ten-construction-seed spectral sweep uses six overlap settings and four budgets, with 960 configurations per sketch rank. At ranks 10, 20, 50, relative exact merged-effective-rank gains over rank-only are respectively 2.09, 3.81, 4.94% for the Gram sketch and 4.90, 5.48, 5.68% for DPP. These gains are in the geometry objective, not classification accuracy. Its entropy-continuity bounds follow the type of bound in [Audenaert \(2007\)](#); all 2,880 selected-prefix intervals cover, but none of 7,200 evaluated greedy steps is certified. The archived source is `results/two_stage_controlled_rank_sweep_corrected_v3/`. These observations do not imply that every low-rank sketch must fail.

The separate E5 study compares sketch payloads to packed symmetric float64 Grams, not inflated dense storage. Among full-grid strategies whose mean byte ratio is at most 25%, the best selected-set match to Full-Gram greedy is 15.42%, achieved by adaptive rank cap 2 with mean byte ratio 14.91%. Its pairwise certificate rate is 3.91% and fully certified selection rate zero. Adaptive fallback reaches full selected-set agreement but costs approximately 160% of packed-Gram bytes on average when previous transmissions are counted. This fails the joint decision-and-communication requirement. The reference is Full-Gram greedy, not the exhaustive global minimizer.

Table 4: Brier top-ten recall (%) is reported for all follow-up scoring rules at $s = 227$. Averages use the same domain-level procedure as Table 3; original-seed controls are excluded.

Screen	Ridge	Logistic
Target A-opt	98.33	98.33
Second-moment MMD	99.38	98.96
Bayesian D-opt	82.29	86.46
Merged effective rank	80.42	85.42
Random (20 repeats)	50.12	48.97
Target energy	41.04	33.96
DPP subspace	28.75	35.42
Domain balance	26.46	31.67

An earlier, different experiment trains source-supervised bottleneck adapters on 21 sources and audits seven targets with four encoders and five adapter seeds. Target-train validation uses three stratified five-fold partitions, with held-out target test accessed separately after selection freezing. Frozen identity reaches mean accuracy 0.83654, an untrained random adapter 0.77144, and the validation-oracle source-supervised adapter 0.76951. The last trails identity by 0.06703, with target-cluster 95% interval $[-0.09771, -0.04400]$. Small shortlist regret relative to this adapter family is therefore not downstream benefit. This is not a DomainNet experiment or evidence against all adaptation methods. Results reside in `results/two_stage_repeated_cv_utility_corrected_v3/`.

C.1 DECISION-BUDGET REANALYSIS

The post-hoc budget audit uses the four new projection seeds and both readouts without selecting a favorable seed. For each task and endpoint, exhaustive validation chooses its combination using validation losses only. Its test loss is then the paired reference for the frozen shortlist choices. In contrast, the empirical test oracle remains a lower audit reference. All five original budgets are retained.

Table 5: Final logistic-regression Brier and the rate of exceeding the exhaustive-validation reference by more than 0.005 are reported. Means average random repetitions before encoders and seeds within each domain, then the six domains. The tolerance is exploratory, not a new success gate.

Screen	Fits	Test Brier	Gap > 0.005 (%)
Random	227	0.866030	15.52
Target A-opt	50	0.864511	2.08
MMD	50	0.864365	2.08
Exhaustive validation	455	0.864118	0.00

We report absolute-gap sensitivity grids of 0.001, 0.005, 0.01 for Brier, 0.01, 0.05, 0.1 for NLL, and 0.5, 1, 2 percentage points for error. Exceedance indicators are computed before averaging random repetitions; thresholding an averaged loss would answer a different question. Descriptive paired bootstrap ranges use 2,000 resamples of the six domain averages. They do not account for dependence induced by shared source pools, validation resampling, or independent datasets. All endpoint curves, paired domain differences, and threshold rates are in `results/target_conditioned/revision_20260922/budget_audit/`. The script `scripts/analyze_e2_b_decision_budget.py` verifies recorded loss hashes and constructs the exhaustive reference from validation files.




Figure 3: The panels report post-hoc diagnostics for different candidate families. The left panel shows oracle Brier improvement over DPP, averaging encoders within each PACS/Office-Home target. The right panel shows relative best-to-worst Brier spans for the original DomainNet tasks. The original evaluation thresholds are unchanged.

C.2 LOSS RESOLUTION AND TEST-SAMPLE STABILITY

The post-hoc loss-span diagnostic in Appendix Figure 3 puts the original regret criterion in context. Every candidate’s relative Brier regret lies below the 5% threshold, so that threshold cannot distinguish successful screens in this experiment. Recall provides the more informative comparison, but it raises another question: how sensitive is the test top ten to the sampled test examples?

A CUDA reconstruction first reproduces all original Brier top-ten sets. We then apply paired test-example bootstrap resampling while fixing fitted models and validation choices. Although resampled and original top-ten sets overlap by only 85.15% on average, A-opt and MMD retain 97.42% and 97.85% of the resampled top tens. Broad shortlists are less sensitive than exact rank membership to these test-sample changes.

Changing the metric has a larger effect on which combinations count as good: the original Brier and classification-error top tens overlap by only 29.17% on average. A screen should therefore be evaluated against the intended deployment metric, even when its average recall is high for several metrics. The resampling protocol and additional metric comparisons appear in Appendix B.

C.3 CANDIDATE CONSTRUCTION AND DOMAIN CONTROLS

The fixed challenge uses projection seed 20260917, both encoders, and both readouts. Each domain’s original 1,536 source images are sorted by SHA-256 of a fixed seed, domain identifier, and sample identifier, then partitioned into three blocks of 512. No labels, feature values, or outcomes determine membership. Target-selection, validation, and test sets do not change. The original construction is rerun as a control: all 65,520 recorded validation/test loss values agree exactly with the corresponding earlier seed. Repartitioning does not reuse CLIP stratum labels, but still conditions on the originally selected source images; it is neither an independent dataset nor a new test sample.

Mean matching minimizes $\|\bar{h}_S - \bar{h}_T\|_2^2$, the linear-kernel MMD score. Domain matching additionally uses source-domain identifiers: it replaces each block moment by its full source-domain moment, weights these moments by the candidate’s domain sample proportions, and minimizes squared Frobenius distance to the target moment. The original challenge breaks ties lexicographically. Because domain matching creates many structural ties, an additional post-hoc audit randomizes the

order within each tied domain-composition stratum using 20 fixed seeds. It freezes every new validation choice before reading test losses and does not refit or change classifier parameters. The original tie protocol remains available in the released curves.

Table 6: The construction challenge uses logistic Brier. Final loss is validation-selected at $s = 50$. The last columns measure retention of one empirical optimum when half of each non-singleton domain-composition stratum is kept, not final selection accuracy. Means average strata and encoders within targets, then six targets. Domain matching uses 20 randomized tie orders for final loss; its within-stratum entry is the exact uniform-tie expectation.

Screen	Final Brier ($s = 50$)		Within-composition (%)	
	Original	Hash	Original	Hash
Target A-opt	0.863170	0.869699	77.22	55.83
MMD	0.863110	0.869257	76.11	55.00
Mean matching	0.863713	0.873038	75.28	57.78
Domain moment (random ties)	0.863667	0.869356	45.68	45.68
Random	0.870013	0.874244	45.04	45.31

There are 20 strata of nine candidates and ten of 27 per target; five singleton strata are excluded. Within a stratum, every candidate uses the same domain multiplicities and total sample count. Random repeats are averaged first, followed by strata, encoders, and target domains. A domain-only score is constant inside a stratum, giving expected optimum retention $(20 \cdot 4/9 + 10 \cdot 13/27)/30 = 45.68\%$ under uniform ties. Lexicographic within-stratum results are recorded but not interpreted as domain discrimination.

After hash repartitioning, logistic classification-error optimum retention is 41.94% for A-opt and 35.28% for MMD, versus 45.75% for random selection. Mean Brier omission remains lower than random for both scores, but the advantage is smaller than on the original blocks. These metric-dependent observations describe conditional ranking performance; they do not establish a population guarantee or isolate the causal contribution of heterogeneity from the properties of this particular partition.

The complete curves and domain-conditional audits are in `results/target_conditioned/revision_20260922/challenge_summary/`; randomized domain ties are in `results/target_conditioned/revision_20260922/domain_tie_audit/`. The raw three-stage records bind candidate membership, source features, configuration, upstream choices, and numerical payloads by hashes. Local summaries verify all numerical hashes and upstream seals; frozen model hashes were checked on the server before test access.

C.4 MATCHED QUALITY AND COMPUTATION COSTS

The timing study uses the original candidate construction and projection seed 20260917. Each encoder-target-readout task runs A-opt, MMD, random screening, and exhaustive validation independently, with three repetitions and a seeded method order. BLAS uses one thread. Shortlist checkpoints measure cumulative elapsed time after a single ranking; no fitted model is shared across methods. Ridge lazily caches each used block’s float64 Gram and label right-hand side, including in exhaustive search. Preliminary float32-Gram ridge timings are retained in the archive but excluded from all tables.

Every timed validation loss and choice is checked against the frozen scientific audit. Quality in Figure 4 uses the actual choices from these same three runs, with Brier, NLL, and error each using its matching validation metric. Thus the random quality coordinates here average three timed rankings, not the 20 rankings in the main scientific audit. We first average repetitions within each task, then encoders within targets and the six targets equally; timing medians are retained as diagnostics.

The resident scope excludes loading stage bundles, checking their integrity, and projecting cached features. Table 7 accounts for these operations separately for two six-target encoder batches. Target-aware methods load both selection and validation bundles; random and exhaustive methods need only validation bundles. Shared preparation is charged once per encoder batch, not converted into a fictitious single-task time.




Figure 4: Final test Brier is plotted against matched decision time and additive image-cost accounting. Each curve has budgets 10, 25, 50, 100, 227; crosses mark exhaustive validation over 455 candidates. Rows have different vertical scales. Means use two encoders, six targets, and three independent executions per method. Raw-image costs exclude pool construction, downloads, cache serialization, and test auditing.

Table 7: Disk-cache batch costs in seconds include input loading, integrity checks, projection, and all twelve encoder–target decisions. They are totals for two six-target batches, whereas Table 1 reports per-task means. The choices and losses do not change.

Readout	A-opt (50)	MMD (50)	Random (227)	Full (455)
Ridge	9.35	5.46	10.40	19.18
Logistic	56.77	54.63	255.89	515.12

Raw-image profiling measures model initialization, image reads and hash checks, decoding, preprocessing, inference, CPU feature transfer, and normalization/projection. Each source block is timed separately; random screening is charged only for the union of blocks in its actual prefix, while geometric scoring and exhaustive validation use all 15 eligible source blocks. Each method also needs target-validation features; only target-aware scores need the target-selection sample. No test images are processed. The fixed candidate inventory is supplied to this measurement, so CLIP-based pool construction is outside its scope.

The machine uses an Intel Xeon Platinum 8352V CPU and an NVIDIA vGPU-32GB. Feature extraction uses CUDA 12.4, PyTorch 2.5.1, float32 inference without TF32, batch size 64, and four data-loading workers. The readout stack uses NumPy 1.26.4, SciPy 1.16.1, and scikit-learn 1.7.1. Feature components are measured in one pass, with operating-system disk caches left unchanged. Block-level loader startup is included. The totals therefore describe this measured implementation and reuse scenario, not a hardware-independent speedup.

All 32 quality/time curve entries, three loss endpoints, paired differences, and 384 task-level budget entries are in `results/target_conditioned/revision_20260922/cost_summary/`. Standalone timing runs and per-block feature profiles accompany them. The manuscript verifier independently reconstructs the matched quality coordinates, random block unions, and all additive totals from those records.